Machine Learning On Edge Computing

Presented by:
Rohit Singh
Akanksha Raina
Srujana Malisetti

Contents


Motivation


Problem Statement


Goal

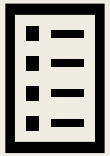
Related Work


Design


Measurement


Expected Result


Timeline

Machine Learning Services

- Many Cloud Providers now a days are providing Machine Learning Services.
- Termed as MLaaS.




Status Quo Approach

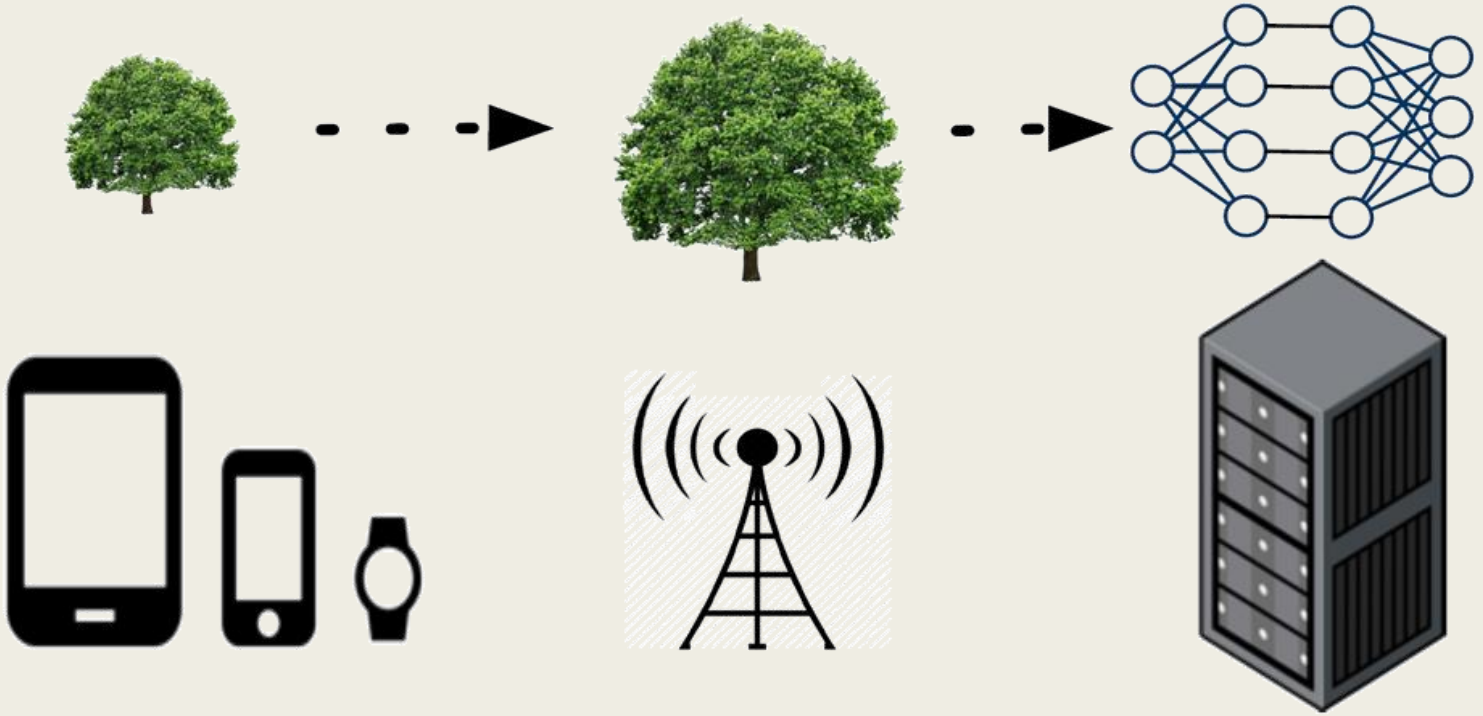

On Other side

- Intelligent Personal Assistants running on SoC integration devices, have capability to run ML Models efficiently.





How about Edge Computing ?

Many Options

?

AlexNet VGG CaffeNet	Image Classification
DeepFace FaceNet NormFace	Face Recognition
Kaldi DeepSpeech	Speech Recognition
SENNA Tesseract	Text Recognition

N

x

Apple Siri
Microsoft Cortana
Google Now
Amazon Alexa
Raspberry Pi
Jetson Nano
Cloud - VM, Container, Functions

M

Choose Best?

N models. → M devices. → N x M possibilities.

How to choose best devices or models?

Help from!!

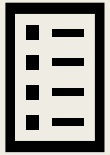

- **Complexity v/s Performance : Empirical Analysis of Machine Learning as a Service**

http://people.cs.uchicago.edu/~ravenben/publications/pdf/mlaas-imc17.pdf

- **Neurosurgeon: Collaborative Intelligence Between the Cloud and Mobile Edge**

http://web.eecs.umich.edu/~jahausw/publications/kang2017neurosurgeon.pdf

- **Spock: Exploiting Serverless Functions for SLO and Cost Aware Resource Procurement in Public Cloud**

http://www.cse.psu.edu/~pxt176/publications/cloud-spock.pdf

- **Distributed Perception by Collaborative Robots**

https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=8411096

Complexity v/s Performance : Empirical Analysis of Machine Learning as a Service

- The paper discusses how MLaaS systems can provide an alternative to standalone ML classifiers.
- The paper provides empirical analysis of MlaaS platforms. Following points were observed during the analysis:
  - With more control comes more potential performance gains as well as greater performance degradation from poor configuration decisions.
  - Fully automated platforms are optimizing classifiers using internal tests.
  - Much of the gains from configuration and tuning come from choosing the right classifier.
  - Experimenting with a small random subset of classifiers is likely to achieve near optimal results.

Neurosurgeon: Collaborative Intelligence Between the Cloud and Mobile Edge

- A system that can automatically partition Deep Neural Networks between mobile devices and the cloud at the granularity of neural network layers.
- Neurosurgeon adapts to dynamic conditions, like server load levels and wireless network connection., while making a decision.
- It chooses partition point for best latency and best mobile energy consumption.

Across 8 benchmarks	Average	Maximum
Latency	3.1x	40.7x
Mobile energy Consumption	59.5%	94.7%
Datacenter Throughput	1.5x	6.7x

Spock: Exploiting Serverless Functions for SLO and Cost Aware Resource Procurement in Public Cloud

- The paper describes using serverless functions for resource procurement in public cloud of VM –based autoscaling.
- Spock, a new scalable and elastic control system that exploits both VMs and serverless functions to reduce cost and ensure SLO for elastic web services.
- Spock helps in overcoming the shortcomings of VM-based resource procurement.

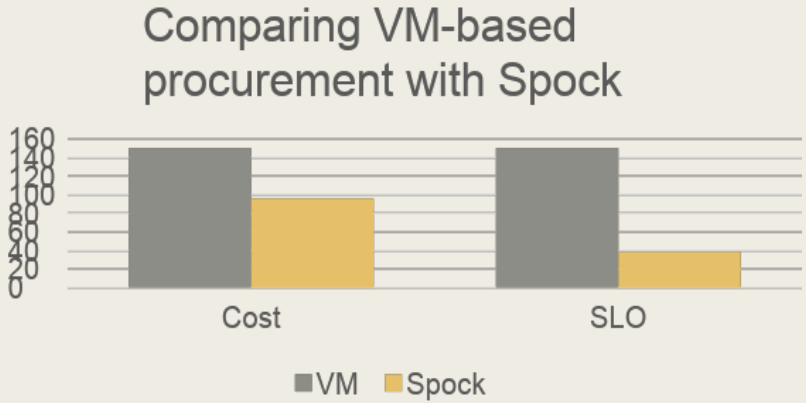

Distributed Perception by Collaborative Robots

- The paper introduces the concept of collaborative approach among robots.
- It enables efficient, dynamic and real time recognition.
- Similar performance results when compared to High Performance machine (HPC) and Jetson TX2

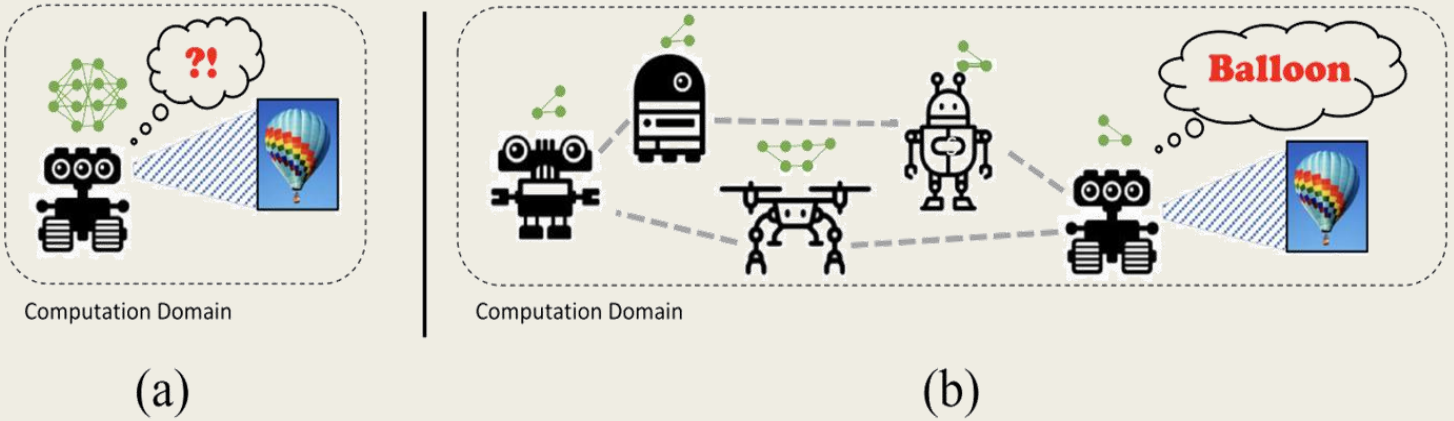

Design

- **List of Devices** :
  - *Raspberry Pi*
  - *Mobile Device (Android)*
  - *AWS Cloud - EC2 Instance, Container, Lambda Functions*
  - *Jetson Nano*
- **List of Machine learning frameworks:**
  - *Caffe*
  - *TensorFlow*
  - *mxNet*
  - *Paddle*
- **List of ML Models or Applications:**
  - *AlexNet*
  - *GoogleNet*
  - *CaffeNet*
  - *DeepFace*
  - *VGG*
  - *SENNA*

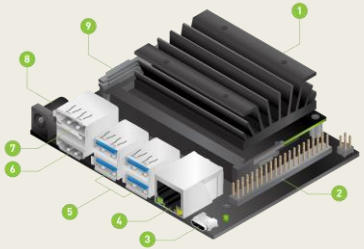

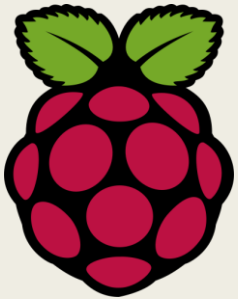

Approach


Deploy frameworks on selected devices
Deploy Models on the devices
Run Experiments
Analyze the results
Generate Output

Measurement

- Baseline - Neurosurgeon
- Performance – Accuracy of calculations (F1-Score, GFLOPS)
- Latency of each model on each device
- Memory consumption on each device
- Battery consumption on each device

Expecting Result

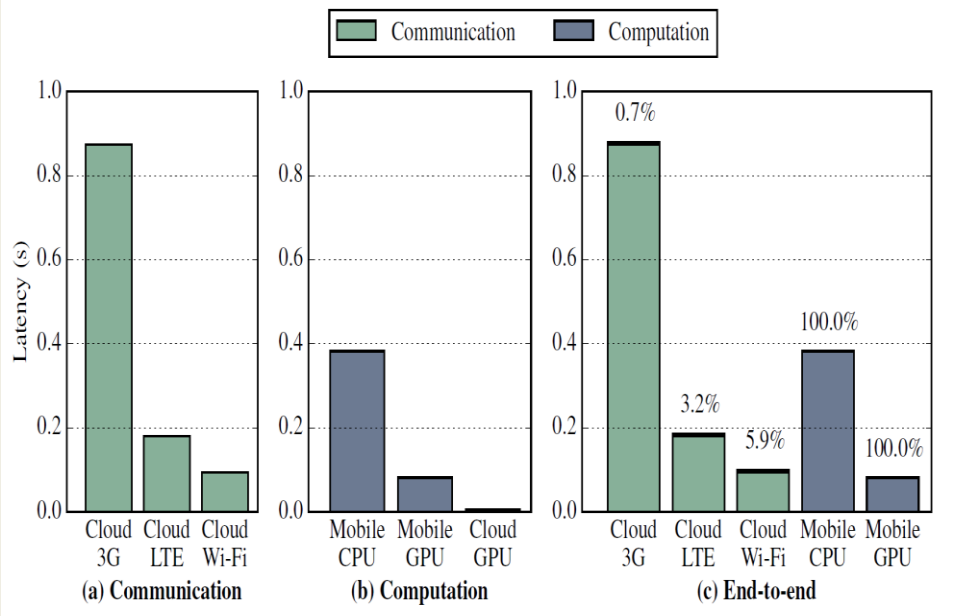

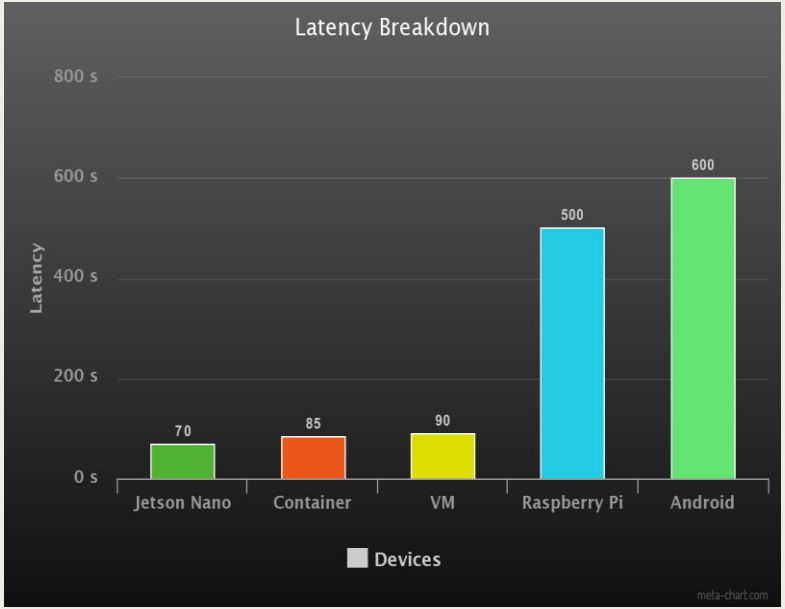

Timeline


Research & Learn
Experiment
Report
23 Sep. – 6 Oct.
7–27 Oct.
28 Oct. – 10 Nov.
11–24 Nov.
25 Nov. – 1 Dec.
Set Up Devices
Analyze

References

- https://skymind.ai/wiki/comparison-frameworks-dl4j-tensorflow-pytorch#ml
- https://medium.com/coinmonks/paper-review-of-alexnet-caffenet-winner-in-ilsvrc-2012-image-classification-b93598314160
- https://arxiv.org/pdf/1804.06655.pdf
- https://en.wikipedia.org/wiki/Speech_recognition
- https://github.com/PaddlePaddle/Paddle

Thank You!!

Questions??